Data Scientist Role Play: Profiling and Analyzing the Yelp Dataset Coursera Worksheet

This is a 2-part assignment. In the first part, you are asked a series of questions that will help you profile and understand the data just like a data scientist would. For this first part of the assignment, you will be assessed both on the correctness of your findings, as well as the code you used to arrive at your answer. You will be graded on how easy your code is to read, so remember to use proper formatting and comments where necessary.

In the second part of the assignment, you are asked to come up with your own inferences and analysis of the data for a particular research question you want to answer. You will be required to prepare the dataset for the analysis you choose to do. As with the first part, you will be graded, in part, on how easy your code is to read, so use proper formatting and comments to illustrate and communicate your intent as required.

For both parts of this assignment, use this "worksheet." It provides all the questions you are being asked, and your job will be to transfer your answers and SQL coding where indicated into this worksheet so that your peers can review your work. You should be able to use any Text Editor (Windows Notepad, Apple TextEdit, Notepad ++, Sublime Text, etc.) to copy and paste your answers. If you are going to use Word or some other page layout application, just be careful to make sure your answers and code are lined appropriately. In this case, you may want to save as a PDF to ensure your formatting remains intact for you reviewer.

Part 1: Yelp Dataset Profiling and Understanding

1. Profile the data by finding the total number of records for each of the tables below:

Code used:

```
select count(*)
from Attribute

select count(*)
from Business

select count(*)
from Category

select count(*)
from Checkin

select count(*)
from elite_years

select count(*)
from friend

select count(*)
from hours

select count(*)
from photo

select count(*)
from review

select count(*)
from tip

select count(*)
from user
```

Result found:

Attribute table =	*10000*
Business table =	*10000*
Category table =	*10000*
Checkin table =	*10000*
elite_years table =	*10000*
friend table =	*10000*
hours table =	*10000*
photo table =	*10000*
review table =	*10000*
tip table =	*10000*
user table =	*10000*

2. Find the total distinct records by either the foreign key or primary key for each table. If two foreign keys are listed in the table, please specify which foreign key.

Code used:

```
select count(distinct id)
from business

select count(distinct business_id)
from hours

select count(distinct business_id)
from category

select count(distinct business_id)
from attribute

select count(distinct id), count(distinct business_id), count(distinct user_id)
from review

select count(distinct business_id)
from checkin

select count(distinct id), count(distinct business_id)
from photo

select count(distinct user_id), count(distinct business_id)
from tip

select count(distinct id)
from user

select count(distinct user_id)
from friend

select count(distinct user_id)
from elite_years
```

Result found:

Business =	*10000*	*(id)*		
Hours =	*1562*	*(business_id)*		
Category =	*2643*	*(business_id)*		
Attribute =	*1115*	*(business_id)*		
Review =	*10000*	*(id- primary),*	*8090 (business_id- foreign),*	*9581 (user_id- foreign)*
Checkin =	*493*	*(business_id)*		
Photo =	*10000*	*(id- primary),*	*6493 (business_id- foreign)*	
Tip =	*537*	*(user_id- foreign),*	*3979 (business_id- foreign)*	
User =	*10000*	*(id)*		
Friend =	*11*	*(user_id)*		
Elite_years =	*2780*	*(user_id)*		

Note: Primary Keys are denoted in the ER-Diagram with a yellow key icon.

3. Are there any columns with null values in the Users table? Indicate "yes," or "no."

Answer: No

SQL code used to arrive at answer:

```
select *
from user
where id is null
or name is null
or review_count is null
or yelping_since is null
or useful is null
or funny is null
or cool is null
or average_stars is null
or compliment_hot is null
or compliment_more is null
or compliment_profile is null
or compliment_cute is null
or compliment_list is null
or compliment_note is null
or compliment_plain is null
or compliment_cool is null
or compliment_funny is null
or compliment_writer is null
or compliment_photos is null
```

4. For each table and column listed below, display the smallest (minimum), largest (maximum), and average (mean) value for the following fields:

i. Table: Review, Column: Stars
min: 1 *max: 5* *avg: 3.7082*

Code used:

```
select min(stars) as min, max(stars) as max, avg(stars) as avg
from review
```

Result found:

```
+-----+-----+--------+
| min | max |    avg |
+-----+-----+--------+
|   1 |   5 | 3.7082 |
+-----+-----+--------+
```

ii. Table: Business, Column: Stars

min: 1.0 *max: 5.0* *avg: 3.6549*

```
select min(stars) as min, max(stars) as max, avg(stars) as avg
from business
```

```
+-----+-----+--------+
| min | max |    avg |
+-----+-----+--------+
| 1.0 | 5.0 | 3.6549 |
+-----+-----+--------+
```

iii. Table: Tip, Column: Likes

min: 0 *max: 2* *avg: 0.0144*

```
select min(likes) as min, max(likes) as max, avg(likes) as avg
from tip
```

```
+-----+-----+--------+
| min | max |    avg |
+-----+-----+--------+
|   0 |   2 | 0.0144 |
+-----+-----+--------+
```

iv. Table: Checkin, Column: Count

min: 1 *max: 53* *avg: 1.9414*

```
select min(count) as min, max(count) as max, avg(count) as avg
from checkin
```

```
+-----+-----+--------+
| min | max |    avg |
+-----+-----+--------+
|   1 |  53 | 1.9414 |
+-----+-----+--------+
```

v. Table: User, Column: Review_count

min: 0 *max: 2000* *avg: 24.2995*

```
select min(review_count) as min, max(review_count) as max, avg(review_count) as avg
from user
```

```
+-----+------+---------+
| min |  max |     avg |
+-----+------+---------+
|   0 | 2000 | 24.2995 |
+-----+------+---------+
```

5. List the cities with the most reviews in descending order:
SQL code used to arrive at answer:

```
select city, sum(review_count) as total_reviews
from business
group by 1
order by 2 desc
```

Copy and Paste the Result Below:

city	total_reviews
Las Vegas	82854
Phoenix	34503
Toronto	24113
Scottsdale	20614
Charlotte	12523
Henderson	10871
Tempe	10504
Pittsburgh	9798
Montréal	9448
Chandler	8112
Mesa	6875
Gilbert	6380
Cleveland	5593
Madison	5265
Glendale	4406
Mississauga	3814
Edinburgh	2792
Peoria	2624
North Las Vegas	2438
Markham	2352
Champaign	2029
Stuttgart	1849
Surprise	1520
Lakewood	1465
Goodyear	1155

(Output limit exceeded, 25 of 362 total rows shown)

6. Find the distribution of star ratings to the business in the following cities:
i. Avon
SQL code used to arrive at answer:

```
select stars as star_rating, count(stars) as count
from business
where city = 'Avon'
group by 1
order by 1 desc
```

Copy and Paste the Resulting Table Below (2 columns – star rating and count):

star_rating	count
5.0	1
4.5	1
4.0	2
3.5	3
2.5	2
1.5	1

ii. Beachwood
SQL code used to arrive at answer:

```
select stars as star_rating, count(stars) as count
from business
where city = 'Beachwood'
group by 1
order by 1 desc
```

Copy and Paste the Resulting Table Below (2 columns – star rating and count):

star_rating	count
5.0	5
4.5	2
4.0	1
3.5	2
3.0	2
2.5	1
2.0	1

7. Find the top 3 users based on their total number of reviews:
SQL code used to arrive at answer:

```
select id, name, sum(review_count) as 'review count'
from user
group by 1
order by 3 desc
limit 3
```

Copy and Paste the Result Below:

id	name	review count
-G7Zkl1wIWBBmD0KRy_sCw	Gerald	2000
-3s52C4zL_DHRK0ULG6qtg	Sara	1629
-8lbUNlXVSoXqaRRiHiSNg	Yuri	1339

8. Does posing more reviews correlate with more fans?

Please explain your findings and interpretation of the results: *Not really. I have written a query to find out the 5 top most review writers among the users. From the output, the user named Sara would have more fans than Yuri, .Hon and William as she wrote more reviews than they did but that didn't happen. So, apparently more reviews don't directly correlate with more fans.*

```
select id, name, sum(review_count), sum(fans)
from user
group by 1
order by 3 desc
limit 5
```

id	name	review_count	fans
-G7Zkl1wIWBBmD0KRy_sCw	Gerald	2000	253
-3s52C4zL_DHRK0ULG6qtg	Sara	1629	50
-8lbUNlXVSoXqaRRiHiSNg	Yuri	1339	76
-K2Tcgh2EKX6e6HqqIrBIQ	.Hon	1246	101
-FZBTkAZEXoP7CYvRV2ZwQ	William	1215	126

9. Are there more reviews with the word "love" or with the word "hate" in them?
Answer: *There are more reviews with the word 'Love' than 'Hate'. Sounds good!*
SQL code used to arrive at answer:

```
select 'Love' Word, count (text) as 'Count'
from review
where text like '%love%'
union
select 'Hate' Word, count (text) as 'Count'
from review
where text like '%hate%'
```

Word	Count
Hate	232
Love	1780

10. Find the top 10 users with the most fans:
SQL code used to arrive at answer:

```
select name, sum(fans) as total_fans
from user
group by 1
order by 2 desc
limit 10
```

Copy and Paste the Result Below:

name	total_fans
Amy	519
Mimi	498
Harald	311
Gerald	256
Lisa	207
Nicole	200
Christine	187
Mark	156
Jen	148
Linda	148

11. Is there a strong relationship (or correlation) between having a high number of fans and being listed as "useful" or "funny?" Out of the top 10 users with the highest number of fans, what percent are also listed as "useful" or "funny"?

Key:
0% - 25% - Low relationship
26% - 75% - Medium relationship
76% - 100% - Strong relationship

SQL code used to arrive at answer:

Query 1: the following query finds out the total number of fans, total number of funny reviews and total number of useful reviews.

```
select sum(fans) as 'total fans', sum(funny) as 'total funny', sum(useful) as 'total
useful'
from user;
```

Result:

total fans	total funny	total useful
14896	247927	380563

Query 2: The following query finds out the total number of fans, total number of funny reviews and total number of useful reviews for the top 10 users with the most fans. It also finds out the percentage in the following columns. For example, Amy has 3.48% fans of the total fans, 1.04% funny reviews/reactions of the total funny reviews and 0.88% useful reviews among the total reviews.

```
select name, sum(fans), sum(funny), sum(useful),
  round((sum(fans)/14896.)*100, 2)  as '%_fans',
  round((sum(funny)/247927.)*100, 2)  as '%_funny',
  round((sum(useful)/380563.)*100, 2)  as '%_useful'
from user
group by 1
order by 2 desc
limit 10
```

name	sum(fans)	sum(funny)	sum(useful)	%_fans	%_funny	%_useful
Amy	519	2578	3366	3.48	1.04	0.88
Mimi	498	143	266	3.34	0.06	0.07
Harald	311	122419	122921	2.09	49.38	32.3
Gerald	256	2326	17530	1.72	0.94	4.61
Lisa	207	144	533	1.39	0.06	0.14
Nicole	200	187	944	1.34	0.08	0.25
Christine	187	6672	5305	1.26	2.69	1.39
Mark	156	703	4250	1.05	0.28	1.12
Jen	148	3426	4029	0.99	1.38	1.06
Linda	148	3080	3946	0.99	1.24	1.04

Query 3: the following query finds out how much the top 10 users' fans constitute of the total fans. It also shows the percentage of funny and useful reviews. The last two columns show the relationship between fans and funny in percentage (27.35%) and the relationship between fans and useful in percentage (41.98%).

```
select
  round( sum(fans/14896.*100), 2) as 'top 10 fan total (%)',
  round(sum(funny/247927.*100), 2) as 'top 10 funny total (%)',
  round(sum(useful/380563.*100), 2) as 'top 10 useful total (%)',
  round((sum(fans/14896.*100))*100/(sum(funny/247927.*100)), 2) as '%_fantofunny',
  round((sum(fans/14896.*100))*100/(sum(funny/380563.*100)), 2) as '%_fantouseful'
```

```
from user
where name in ('Amy', 'Mimi', 'Harald', 'Gerald', 'Christine', 'Lisa', 'Cat', 'William',
'Fran', 'Lissa')
order by 2 desc;
```

top 10 fan total (%)	top 10 funny total (%)	top 10 useful total (%)	%_fantofunny	%_fantouseful
16.78	61.34	44.86	27.35	41.98

Conclusion: Relationship between the high number of fans and being listed "funny" = 27.35% Relationship between the high number of fans and being listed "useful" = 41.98%. So, there's a Medium relationship (26% - 75%) between these factors.

Part 2: Inferences and Analysis

1. Pick one city and category of your choice and group the businesses in that city or category by their overall star rating. Compare the businesses with 2-3 stars to the businesses with 4-5 stars and answer the following questions. Include your code.

   City = Las Vegas
   Category = Restaurants

name	star_rating	review_count	hours	days	postal_code	latitude	longitude
Big Wong Restaurant	4-5 stars	768	Friday|10:00-23:00	weekdays	89146	36.1267	-115.21
Big Wong Restaurant	4-5 stars	768	Monday|10:00-23:00	weekdays	89146	36.1267	-115.21
Big Wong Restaurant	4-5 stars	768	Saturday|10:00-23:00	weekends	89146	36.1267	-115.21
Big Wong Restaurant	4-5 stars	768	Sunday|10:00-23:00	weekends	89146	36.1267	-115.21
Big Wong Restaurant	4-5 stars	768	Thursday|10:00-23:00	weekdays	89146	36.1267	-115.21
Big Wong Restaurant	4-5 stars	768	Tuesday|10:00-23:00	weekdays	89146	36.1267	-115.21
Big Wong Restaurant	4-5 stars	768	Wednesday|10:00-23:00	weekdays	89146	36.1267	-115.21
Jacques Cafe	4-5 stars	168	Friday|11:00-20:00	weekdays	89134	36.1933	-115.304
Jacques Cafe	4-5 stars	168	Monday|11:00-20:00	weekdays	89134	36.1933	-115.304
Jacques Cafe	4-5 stars	168	Saturday|11:00-20:00	weekends	89134	36.1933	-115.304
Jacques Cafe	4-5 stars	168	Sunday|8:00-14:00	weekends	89134	36.1933	-115.304
Jacques Cafe	4-5 stars	168	Thursday|11:00-20:00	weekdays	89134	36.1933	-115.304
Jacques Cafe	4-5 stars	168	Tuesday|11:00-20:00	weekdays	89134	36.1933	-115.304
Jacques Cafe	4-5 stars	168	Wednesday|11:00-20:00	weekdays	89134	36.1933	-115.304
Wingstop	2-3 stars	123	Friday|11:00-0:00	weekdays	89103	36.1003	-115.21
Wingstop	2-3 stars	123	Monday|11:00-0:00	weekdays	89103	36.1003	-115.21
Wingstop	2-3 stars	123	Saturday|11:00-0:00	weekends	89103	36.1003	-115.21
Wingstop	2-3 stars	123	Sunday|11:00-0:00	weekends	89103	36.1003	-115.21
Wingstop	2-3 stars	123	Thursday|11:00-0:00	weekdays	89103	36.1003	-115.21
Wingstop	2-3 stars	123	Tuesday|11:00-0:00	weekdays	89103	36.1003	-115.21
Wingstop	2-3 stars	123	Wednesday|11:00-0:00	weekdays	89103	36.1003	-115.21

i. Do the two groups you chose to analyze have a different distribution of hours? *There's no significant difference in the hours between these two groups. But Wingstop (2-3*) stays open until midnight whereas the (4-5*) restaurants close by 11 pm.*

ii. Do the two groups you chose to analyze have a different number of reviews? *Yes, they do. The (4-5*) restaurants have more than 160 reviews whereas the (2-3*) restaurant has 123 reviews.*

iii. Are you able to infer anything from the location data provided between these two groups? *From the latitude and longitude provided, Big Wong Restaurant is located near the strip in Las Vegas which is a reason it's a big hit. The other one, Jacques Café is located near some businesses, library and a high school. It's also surrounded by residential areas.*

SQL code used for analysis:

```
select b.name,
case
when b.stars between 2 and 3 then '2-3 stars'
when b.stars between 4 and 5 then '4-5 stars'
end as star_rating,
b.review_count,
h.hours,
case
when (hours like "%monday%"
or hours like "%tuesday%"
or hours like "%wednesday%"
or hours like "%thursday%"
or hours like "%friday%") then 'weekdays'
when (hours like "%saturday%"
or hours like "%sunday%") then 'weekends'
end as days,
postal_code,
latitude,
longitude
from business b
join hours h
on b.id = h.business_id
join category c
on c.business_id = b.id
where (b.city like 'las vegas'
and c.category like 'restaurants')
and (b.stars between 2 and 3
or b.stars between 4 and 5)
group by hours
order by review_count desc
```

2. Group business based on the ones that are open and the ones that are closed. What differences can you find between the ones that are still open and the ones that are closed? List at least two differences and the SQL code you used to arrive at your answer.

business	is_open	total_stars	avg_stars	total_reviews	avg_reviews	sum(r.useful)
61	0	71	3.65	9217	130.0	69
446	1	565	3.76	175821	311.0	484

i. Difference 1: *Closed ones have fewer stars and less average stars than the open ones.*
ii. Difference 2: *Closed ones have fewer total reviews than the open ones.*

SQL code used for analysis:

```
select count(distinct b.id) as business,
        b.is_open,
        count(r.stars) as total_stars,
        round(avg(r.stars), 2) avg_stars,
        sum(b.review_count) as total_reviews,
        round(avg(b.review_count), 0) as avg_reviews,
        sum(r.useful)
from business b
join review r
on b.id = r.business_id
group by 2
```

3. For this last part of your analysis, you are going to choose the type of analysis you want to conduct on the Yelp dataset and are going to prepare the data for analysis.

Ideas for analysis include: Parsing out keywords and business attributes for sentiment analysis, clustering businesses to find commonalities or anomalies between them, predicting the overall star rating for a business, predicting the number of fans a user will have, and so on. These are just a few examples to get you started, so feel free to be creative and come up with your own problem you want to solve. Provide answers, in-line, to all of the following:

i.Indicate the type of analysis you chose to do:

I am going to answer the following question where Category is in (Asian Fusion, Chinese, Mexican, French, Italian, Indonesian, Korean, Japanese, Indian, Buffet, Coffee & Tea, Barbeque, Chicken Wings, Diners, Fast Food)

Part 1: What are the names, average ratings, average number of reviews and geographic location of the restaurants in these categories?
Part 2: Which of these restaurants are open and which are close?
Part 3: Did the customers feel any of the following emotions during their visit to these restaurants and wrote reviews about them?

Awesome!
Happy!
Delicious!
Great!
Yummy!
Best!

ii. Write 1-2 brief paragraphs on the type of data you will need for your analysis and why you chose that data:

We will need to (inner) join the Business table and Category table for Part 1 and Part 2. For Part 3, we will have to (inner) join Business, Review and Category tables.

iii. Output of your finished dataset:

Part 1:

category	name	avg_rating	avg_reviews	city
Asian Fusion	Big Wong Restaurant	3.5	396.5	Las Vegas
Barbeque	Bootleggers Modern American Smokehouse	3.75	252.5	Phoenix
Chicken Wings	The Erin Mills Pump & Patio	3.0	75.0	Mississauga
Chinese	Club India restaurant	3.13	199.0	Edinburgh
Coffee & Tea	Cabin Fever	3.83	80.0	Toronto
Diners	Market Street Diner	3.75	58.5	Sun Prairie
Fast Food	Del Taco	3.21	26.43	Gilbert
French	Jacques Cafe	4.0	128.5	Las Vegas
Indian	Club India restaurant	3.6	12.6	Edinburgh
Italian	Restaurant Rosalie	3.5	74.0	Montréal
Japanese	Hibachi-San	3.8	30.4	Las Vegas
Korean	Seoul Garden Korean Restaurant	4.25	31.5	Cuyahoga Falls
Mexican	Taqueria Y Cenaduria Culiacan	3.5	46.71	Tolleson

Part 2:

category	name	stars	review_count	city	status
Coffee & Tea	Cabin Fever	4.5	26	Toronto	open
Japanese	Hibachi-San	4.5	3	Las Vegas	closed
Korean	Sushi Osaka	4.5	8	Toronto	open
Asian Fusion	Big Wong Restaurant	4.0	768	Las Vegas	open
Barbeque	Bootleggers Modern American Smokehouse	4.0	431	Phoenix	open
French	Edulis	4.0	89	Toronto	open
Italian	Eklectic Pie - Mesa	4.0	129	Mesa	closed
Mexican	Hermanos Mexican Grill	4.0	69	Mississauga	open
French	Jacques Cafe	4.0	168	Las Vegas	closed
Coffee & Tea	Koko Bakery	4.0	162	Cleveland	open
Mexican	Mama Mia	4.0	8	Toronto	closed
Japanese	Masamune Japanese Restaurant	4.0	61	Mississauga	open
Mexican	Miros Cantina Mexicana	4.0	37	Edinburgh	open
Japanese	Naniwa-Taro	4.0	75	Toronto	open
Indian	Patiala House	4.0	10	Brampton	open
Diners	Rise and Dine Cafe	4.0	30	Chesterland	open
Korean	Seoul Garden Korean Restaurant	4.0	55	Cuyahoga Falls	open
Mexican	Taqueria Y Cenaduria Culiacan	4.0	23	Tolleson	open
Indian	Cafe Tandoor	3.5	32	Aurora	open
Indian	Club India restaurant	3.5	3	Edinburgh	closed
Fast Food	Five Guys	3.5	63	Phoenix	open
Indian	Indian Ocean Restaurant	3.5	3	Inverness	open
Diners	Market Street Diner	3.5	87	Sun Prairie	open
Chinese	Ping's Cafe	3.5	21	Fountain Hills	open
Fast Food	Poutine Lafleur	3.5	11	Verdun	open

Part 3:

category	name	stars	review_count	city	customer_emotion	status
Asian Fusion	Big Wong Restaurant	4.0	768	Las Vegas	delicious!	open
Barbeque	Bootleggers Modern American Smokehouse	4.0	431	Phoenix	best!	open
Indian	Cafe Tandoor	3.5	32	Aurora	great!	open

iv. Provide the SQL code you used to create your final dataset:

Part 1:

```
select c.category,
       b.name,
       round(avg(b.stars), 2) as avg_rating,
       round(avg(b.review_count), 2) as avg_reviews,
       b.city
from business b
join category c
on c.business_id = b.id
where c.category in ("Asian Fusion", "Chinese", "Mexican", "French", "Italian",
"Indonesian", "Korean", "Japanese", "Indian", 'Buffet', 'Coffee & Tea', 'Barbeque',
'Chicken Wings', 'Diners', 'Fast Food')
group by 1
```

Part 2:

```
select c.category,
       b.name,
       b.stars,
       b.review_count,
       b.city,
case
when b.is_open = 1 then 'open'
```

```
when b.is_open = 0 then 'closed'
end as 'status'
from business b
join category c
on c.business_id = b.id
where c.category in ("Asian Fusion", "Chinese", "Mexican", "French", "Italian",
"Indonesian", "Korean", "Japanese", "Indian", 'Buffet', 'Coffee & Tea', 'Barbeque',
'Chicken Wings', 'Diners', 'Fast Food')
group by 2
order by 3 desc
```

Part 3:

```
select c.category,
        b.name,
        b.stars,
        b.review_count,
        b.city,
case
when r.text like '%awesome%' then 'awesome!'
when r.text like '%pleased%'
    or r.text like '%happy%' then 'happy!'
when r.text like '%delicious%' then 'delicious!'
when r.text like '%great%' then 'great!'
when r.text like '%tasty%' then 'yummy!'
when r.text like '%best%' then 'best!'
end as 'customer_emotion',
case
when b.is_open = 1 then 'open'
when b.is_open = 0 then 'closed'
end as 'status'
from business b
join hours h
on b.id = h.business_id
join category c
on c.business_id = b.id
join review r
on b.id  = r.business_id
where c.category in ("Asian Fusion", "Chinese", "Mexican", "French", "Italian",
"Indonesian", "Korean", "Japanese", "Indian", 'Buffet', 'Coffee & Tea', 'Barbeque',
'Chicken Wings', 'Diners', 'Fast Food')
group by 2
order by 3 desc
```